

In [40]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

In [2]:

```
claim=pd.read_csv("C:/claims_management/train.csv")
```

In [3]:

```
print('Proceed', round(claim['target'].value_counts()[1]/len(claim) * 100,2), '% of the dataset')
print('Paper Works Remain of', round(claim['target'].value_counts()[0]/len(claim) * 100,2), '% of the dataset')
```

Proceed 76.12 % of the dataset

Paper Works Remain of 23.88 % of the dataset

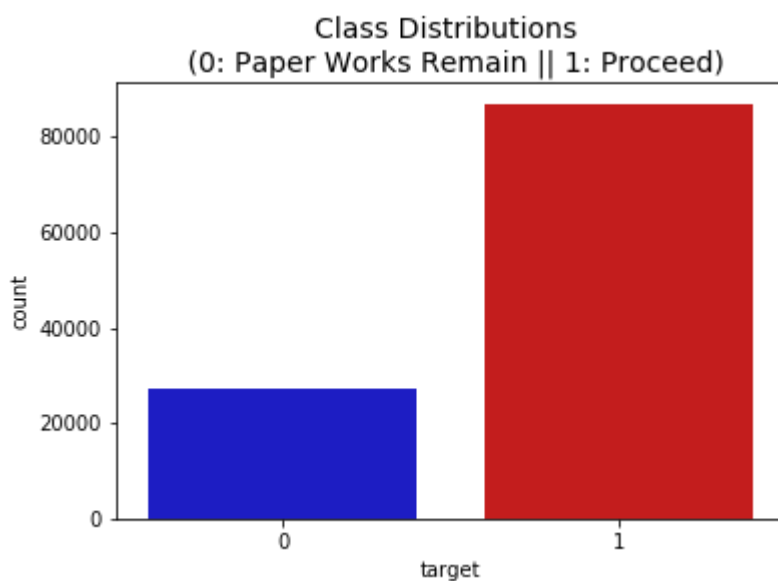
In [4]:

```
colors = ["#0101DF", "#DF0101"]

sns.countplot('target', data=claim, palette=colors)
plt.title('Class Distributions \n (0: Paper Works Remain || 1: Proceed)', fontsize=14)
```

Out[4]:

Text(0.5, 1.0, 'Class Distributions \n (0: Paper Works Remain || 1: Proceed)')



In [5]:

```
def info_null(df):  
    total_null = df.isnull().sum().sort_values(ascending = False)  
    percentage = round(total_null/len(claim) * 100, 2)  
    result = pd.concat([total_null, percentage], axis = 1, keys = ['total Null Values',  
'Percentage'])  
    result.sort_values(by = ['total Null Values', 'Percentage'], ascending = False)  
    return(result[80:100])
```

In [6]:

```
info_null(claim)
```

Out[6]:

	total Null Values	Percentage
v4	49796	43.56
v76	49796	43.56
v2	49796	43.56
v87	48663	42.57
v105	48658	42.56
v98	48654	42.56
v70	48636	42.54
v128	48624	42.53
v5	48624	42.53
v36	48624	42.53
v82	48624	42.53
v81	48624	42.53
v117	48624	42.53
v109	48624	42.53
v108	48624	42.53
v89	48619	42.53
v124	48619	42.53
v25	48619	42.53
v54	48619	42.53
v63	48619	42.53

In [7]:

```
disc=claim.select_dtypes(include = 'object')
```

In [8]:

```

for _ in disc.columns:
    k = {i:d for d, i in enumerate(claim[_].unique())}
    claim[_].replace(k, inplace = True)
for i in claim.columns:
    claim[i].fillna(value = claim[i].mode().iloc[0], inplace = True)
# claim


```

In [9]:

```

from sklearn.feature_selection import SelectKBest, f_classif
X = claim.drop(['ID', 'target'], axis = 1)
y = claim['target']
bestFeature = SelectKBest(score_func = f_classif, k = 20).fit(X, y)
be_features = pd.Series(bestFeature.scores_, X.columns)
best_features = be_features.nlargest(20)
plt.figure(figsize = (20, 10))
best_features.plot(kind = 'bar', color = 'r')
plt.show()

```




In [41]:

```

from sklearn.ensemble import ExtraTreesClassifier
feature = ExtraTreesClassifier().fit(X, y)
feat_importance = pd.Series(feature.feature_importances_, X.columns).nlargest(27)
plt.figure(figsize = (30, 10))
feat_importance.plot(kind = 'bar')
plt.show()

```



In [12]:

```

cormax = claim.corr()
#top_features = cormax.index
# plt.figure(figsize = (200, 200))
# sns.heatmap(cormax, annot=True, cmap="RdYlGn")
# plt.show()


```

In [14]:

```

best_corr_features = cormax['target'].sort_values(ascending = False)[:20][1:]
plt.figure(figsize = (30, 10))
best_corr_features.plot.bar(color = 'g')
plt.show()

```



In [15]:

```
most_imp_features = []
for i in X.columns:
    if i in feat_importance or i in best_corr_features or i in best_features:
        most_imp_features.append(i)
most_imp_features
```

Out[15]:

```
['v2',
 'v4',
 'v10',
 'v12',
 'v14',
 'v21',
 'v22',
 'v24',
 'v30',
 'v31',
 'v34',
 'v38',
 'v40',
 'v44',
 'v47',
 'v48',
 'v50',
 'v52',
 'v56',
 'v59',
 'v61',
 'v62',
 'v64',
 'v66',
 'v71',
 'v72',
 'v75',
 'v79',
 'v91',
 'v100',
 'v106',
 'v107',
 'v110',
 'v112',
 'v113',
 'v114',
 'v119',
 'v123',
 'v125',
 'v129']
```

In [16]:

```
X_imp = claim[most_imp_features]
X_imp
```

Out[16]:

	v2	v4	v10	v12	v14	v21	v22	v24	v30	v31
0	8.727474	3.921026	0.503281	6.085711	11.636387	7.730923	0	0	0	0
1	11.506664	4.915740	1.312910	6.507647	11.636386	6.763110	1	0	0	0
2	5.310079	4.410969	0.765864	6.384670	9.603542	5.245035	2	1	1	0
3	8.304757	4.225930	6.542669	9.646653	14.094723	7.517125	3	2	0	1
4	11.506664	4.915740	1.050328	6.320087	10.991098	6.414567	4	1	1	0
...	...	...	...	...	...	...	...	...	...	...
114316	11.506664	4.915740	1.444201	6.368061	11.865255	7.088172	6905	1	1	0
114317	11.506664	4.915740	6.236324	9.443324	14.924483	8.455263	2371	2	1	1
114318	11.506664	4.915740	2.078775	6.698925	12.269012	6.570625	7435	3	0	1
114319	11.506664	4.915740	1.291029	6.692204	12.573678	7.730751	9742	2	2	0
114320	7.932978	4.640085	0.853391	6.306396	11.967826	7.496000	70	0	0	0

114321 rows × 40 columns

In [17]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_imp, y, test_size = 0.2, random_s
tate = 0)
```

In [42]:

```
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression().fit(X_train, y_train)
lr_train_score = lr.score(X_train, y_train)
print('Accuracy on train data', lr_train_score)
lr_score = lr.score(X_test, y_test)
print('Accuracy on test data', lr_score)
```

Accuracy on train data 0.7726229006298111  
Accuracy on test data 0.7752460091843428

In [19]:

```
from sklearn.metrics import confusion_matrix
confusion = confusion_matrix(y_test, lr.predict(X_test))
print('Logistic regression confusion metrics\n', confusion)
```

```
Logistic regression confusion metrics
[[ 524 4928]
 [ 211 17202]]
```

In [20]:

```
from sklearn.metrics import classification_report
print(classification_report(y_test, lr.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.71	0.10	0.17	5452
1	0.78	0.99	0.87	17413
accuracy			0.78	22865
macro avg	0.75	0.54	0.52	22865
weighted avg	0.76	0.78	0.70	22865

In [21]:

```
y_score_lr = lr.decision_function(X_test)
lr_score_list = list(zip(y_test[0:20], y_score_lr[0:20]))
lr_score_list
```

Out[21]:

```
[(1, -0.36129358961326147),
 (0, 1.2250526716794847),
 (1, 0.5718062255583997),
 (0, 0.13000089332807993),
 (1, 1.439069767345494),
 (1, 0.8241956413054017),
 (1, 0.3476949908467173),
 (0, 0.7144136215053187),
 (0, 0.3707391389158977),
 (1, 0.5897840617186382),
 (0, 0.06855227792471873),
 (1, 1.301339982104703),
 (0, 0.4359688031467399),
 (1, 1.3976182956196244),
 (0, 2.1711515028254897),
 (1, 0.7474904079625183),
 (1, 2.3219650923800867),
 (1, 1.4967916400265795),
 (0, -0.03575264845202199),
 (0, 1.135489654792013)]
```

In [22]:

```
from sklearn.metrics import average_precision_score
average_precision = average_precision_score(y_test, y_score_lr)

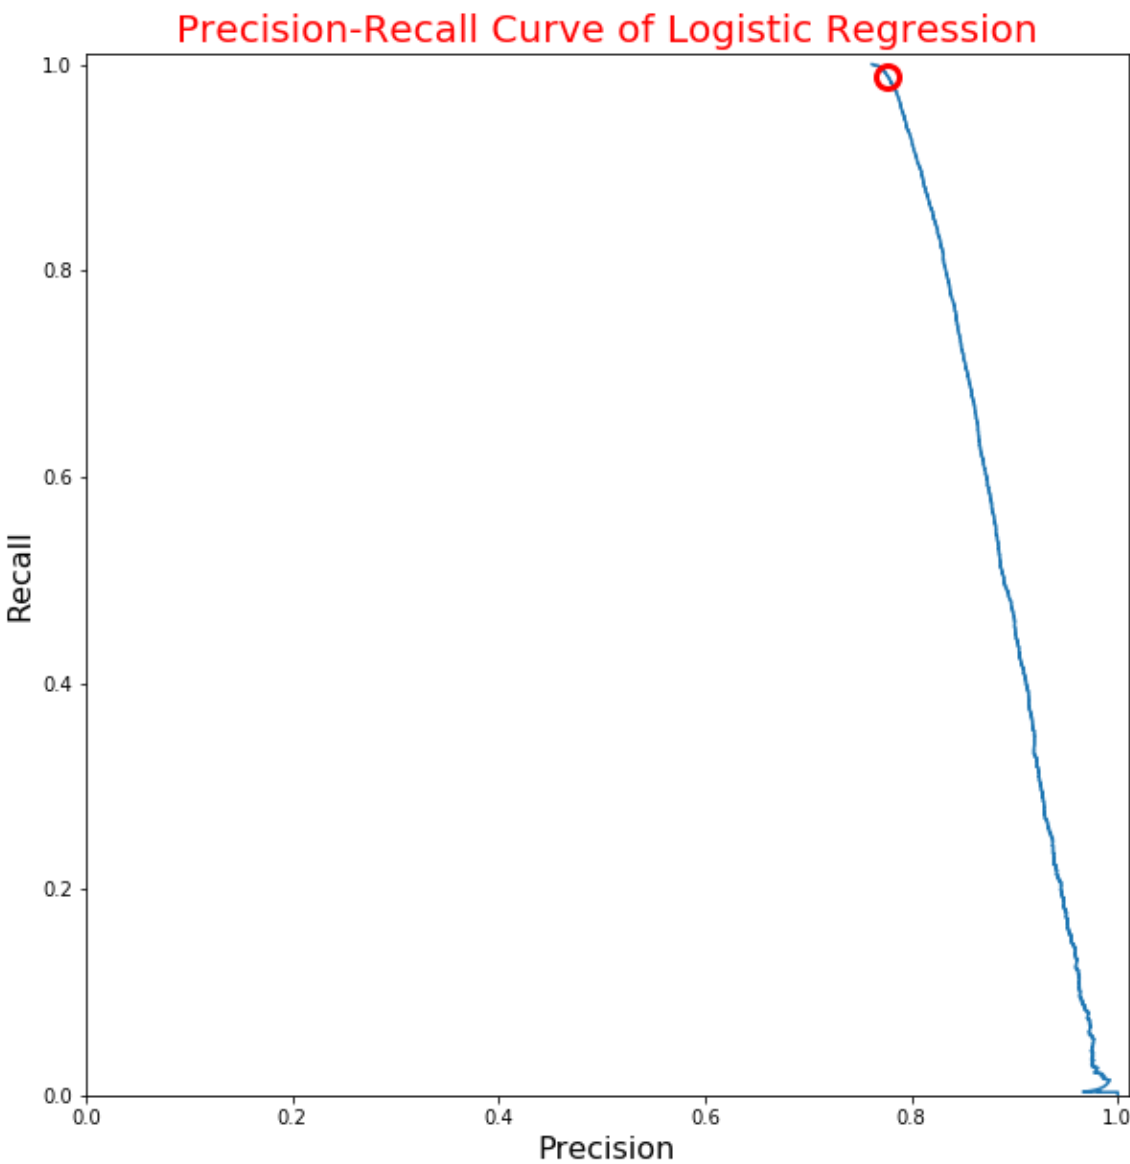
print('Average precision-recall score: {0:0.2f}'.format(
    average_precision))
```

Average precision-recall score: 0.89



In [43]:

```
from sklearn.metrics import precision_recall_curve
precision, recall, thresholds = precision_recall_curve(y_test, y_score_lr)
closest_zero = np.argmin(np.abs(thresholds))
closest_zero_p = precision[closest_zero]
closest_zero_r = recall[closest_zero]
plt.figure(figsize = (10, 10))
plt.xlim([0.0, 1.01])
plt.ylim([0.0, 1.01])
plt.plot(precision, recall, label='Precision-Recall Curve')
plt.plot(closest_zero_p, closest_zero_r, 'o', markersize = 12, fillstyle = 'none', c=
'r', mew=3)
plt.title('Precision-Recall Curve of Logistic Regression', fontsize = 20, c = 'r')
plt.xlabel('Precision', fontsize=16)
plt.ylabel('Recall', fontsize=16)
plt.axes().set_aspect('equal')
plt.show()
```



In [45]:

```
from sklearn.svm import LinearSVC
svc = LinearSVC().fit(X_train, y_train)
print('Accuracy on train data', svc.score(X_train, y_train))
print('Accuracy on test data', svc.score(X_test, y_test))
```

Accuracy on train data 0.7613059832050385

Accuracy on test data 0.7615132298272469

In [25]:

```
from sklearn.metrics import average_precision_score
y_score_svc = svc.decision_function(X_test)

average_precision = average_precision_score(y_test, y_score_svc)


print('Average precision-recall score: {0:0.2f}'.format(
    average_precision))
```

Average precision-recall score: 0.77

In [47]:

```
precision, recall, thresholds = precision_recall_curve(y_test, y_score_svc)
closest_zero = np.argmin(np.abs(thresholds))
closest_zero_p = precision[closest_zero]
closest_zero_r = recall[closest_zero]
plt.figure(figsize = (10, 10))
plt.xlim([0.0, 1.01])
plt.ylim([0.0, 1.01])
plt.plot(precision, recall, label='Precision-Recall Curve')
plt.plot(closest_zero_p, closest_zero_r, 'o', markersize = 12, fillstyle = 'none', c=
'r', mew=3)
plt.title('Precision-Recall Curve of Linear Support Vector Classifier', fontsize = 20,
c = 'b')
plt.xlabel('Precision', fontsize=16)
plt.ylabel('Recall', fontsize=16)
plt.axes().set_aspect('equal')
plt.show()
```

## Precision-Recall Curve of Linear Support Vector Classifier



In [48]:

```
from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier().fit(X_train, y_train)
print('Accuracy on train data', rfc.score(X_train, y_train))
print('Accuracy on test data', rfc.score(X_test, y_test))
```

Accuracy on train data 0.9924225857242828  
Accuracy on test data 0.7509293680297398

In [35]:

```
confusion = confusion_matrix(y_test, lr.predict(X_test))  
print('Random Forest Classifier confusion metrics\n', confusion)
```

Random Forest Classifier confusion metrics

```
[[ 524 4928]  
 [ 211 17202]]
```

In [36]:

```
print(classification_report(y_test, rfc.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.47	0.33	0.39	5452
1	0.81	0.88	0.84	17413
accuracy			0.75	22865
macro avg	0.64	0.61	0.62	22865
weighted avg	0.73	0.75	0.74	22865

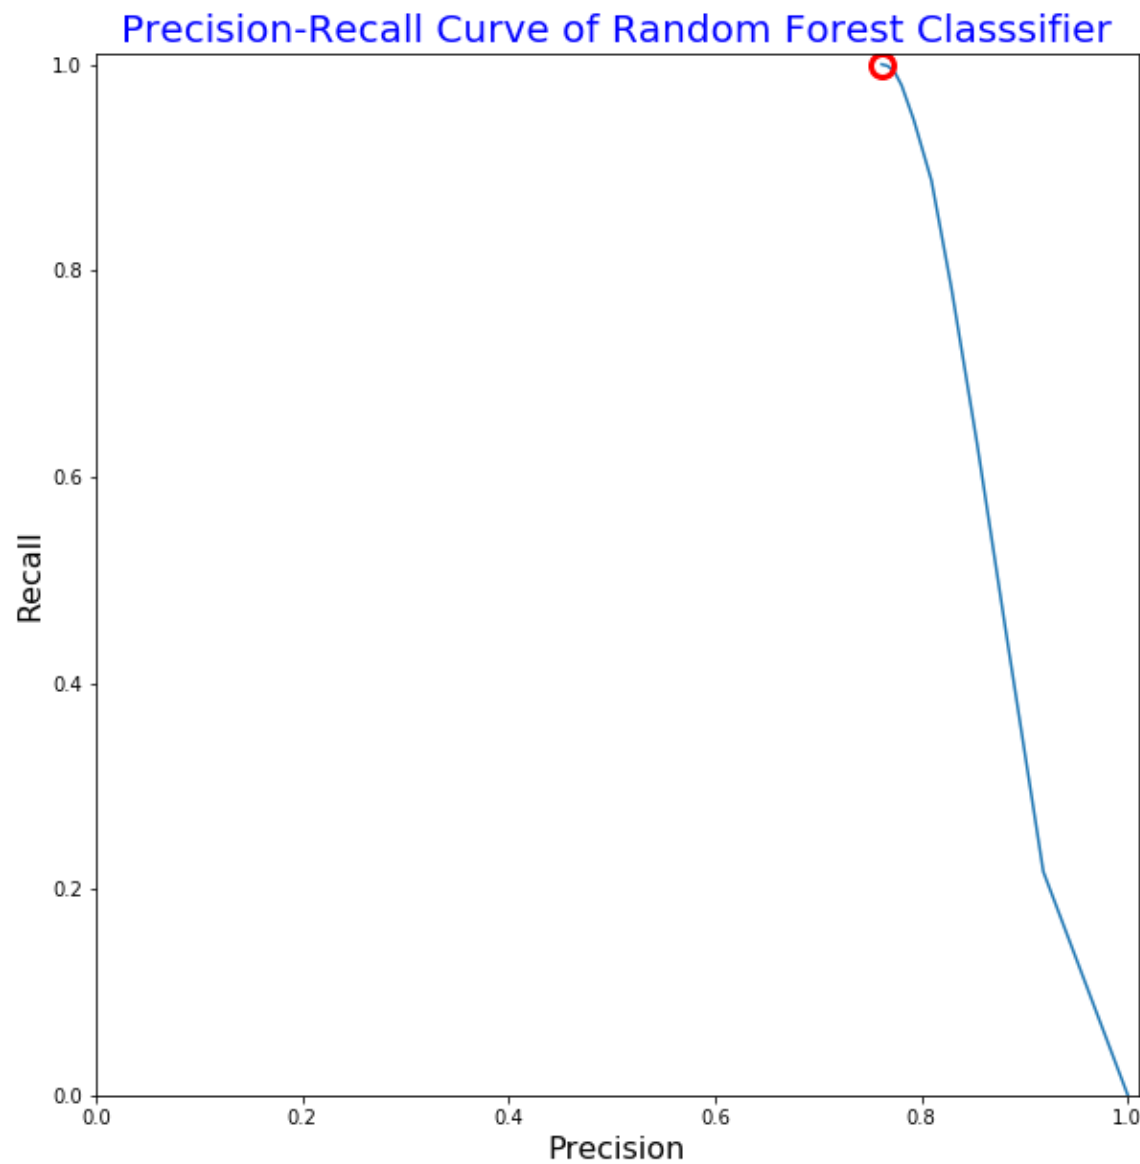
In [28]:

```
y_score_rfc = rfc.predict_proba(X_test)[:,-1]  
print('Average precision-recall score RF: {}'.format(average_precision_score(y_test, y_  
score_rfc)))
```

Average precision-recall score RF: 0.8583035825631781

In [55]:

```
precision, recall, thresholds = precision_recall_curve(y_test, y_score_rfc)
closest_zero = np.argmin(np.abs(thresholds))
closest_zero_p = precision[closest_zero]
closest_zero_r = recall[closest_zero]
plt.figure(figsize = (10, 10))
plt.xlim([0.0, 1.01])
plt.ylim([0.0, 1.01])
plt.plot(precision, recall, label='Precision-Recall Curve')
plt.plot(closest_zero_p, closest_zero_r, 'o', markersize = 12, fillstyle = 'none', c=
'r', mew=3)
plt.title('Precision-Recall Curve of Random Forest Classsifier', fontsize = 20, c = 'b'
)
plt.xlabel('Precision', fontsize=16)
plt.ylabel('Recall', fontsize=16)
plt.axes().set_aspect('equal')
plt.show()
```



In [ ]: